

# CIS 660 DATA MINING

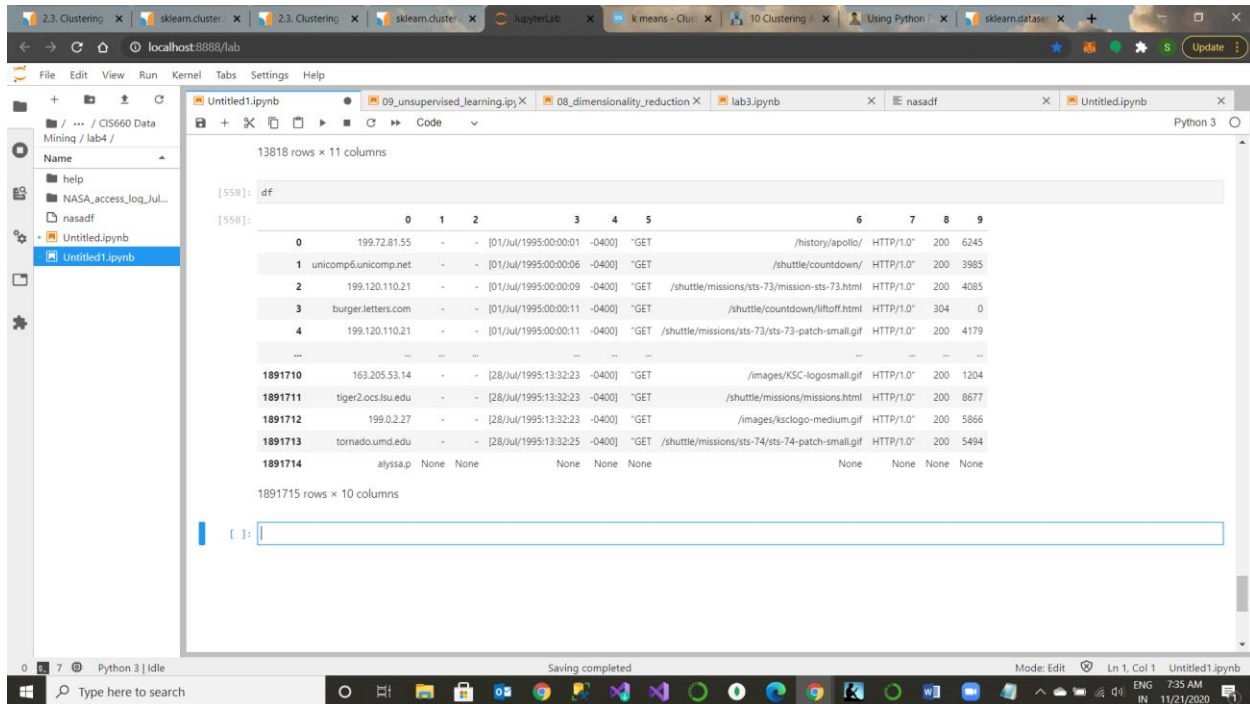
**Lab 4** :Clustering with NASA Webserver Log Data and Feature Selection with PCA tools and also working with different values of K (Extra Credit)

**Fathima Syeda**

**CSU ID-2790024**

## Part 1: Feature Selection, Cleaning, and Preprocessing to Construct an Input from Data Source:

The dataset used in this NASA\_access\_log\_Jul95 server log dataset. This consists of 1800000 logs but of this we a random sample of 14000 logs are chosen .



```
[558]: df
[558]:
```

	0	1	2	3	4	5	6	7	8	9
0	199.72.81.55	-	-	[01/Jul/1995:00:00:01 -0400]	GET	/history/apollo/	HTTP/1.0	200	6245	
1	unicomp6.unicomp.net	-	-	[01/Jul/1995:00:00:06 -0400]	GET	/shuttle/countdown/	HTTP/1.0	200	3985	
2	199.120.110.21	-	-	[01/Jul/1995:00:00:09 -0400]	GET	/shuttle/missions/sts-73/mission-sts-73.html	HTTP/1.0	200	4085	
3	burger.letters.com	-	-	[01/Jul/1995:00:00:11 -0400]	GET	/shuttle/countdown/liftoff.html	HTTP/1.0	304	0	
4	199.120.110.21	-	-	[01/Jul/1995:00:00:11 -0400]	GET	/shuttle/missions/sts-73/sts-73-patch-small.gif	HTTP/1.0	200	4179	
...	...	...	...	...	...	...	...	...	...	...
1891710	163.205.53.14	-	-	[28/Jul/1995:13:32:23 -0400]	GET	/images/KSC-logosmall.gif	HTTP/1.0	200	1204	
1891711	tiger2.ocs.lsu.edu	-	-	[28/Jul/1995:13:32:23 -0400]	GET	/shuttle/missions/missions.html	HTTP/1.0	200	8677	
1891712	199.0.2.27	-	-	[28/Jul/1995:13:32:23 -0400]	GET	/images/ksclogo-medium.gif	HTTP/1.0	200	5866	
1891713	tornado.umd.edu	-	-	[28/Jul/1995:13:32:25 -0400]	GET	/shuttle/missions/sts-74/sts-74-patch-small.gif	HTTP/1.0	200	5494	
1891714	alyssa.p	None	None	None	None	None	None	None	None	None

```
1891715 rows x 10 columns
```

```
[ ]:
```

The dataset consists of 9 columns, namely ---host , 'client\_idntd' , 'user\_id' , 'date\_time' 'method','endpoint' , 'protocol', 'response\_code' , 'content\_size' .

- The client\_idntd and user\_id columns contain null values and hence those columns are dropped .
- The date -time column is transformed into a pandas data time object and additional hour, min, sec columns are found from it and added to the df.
- Certain rows have missing content\_size values and irregular response codes, those rows are eliminated.

After cleaning the dataset looks as follows:

The screenshot shows a JupyterLab interface with a file explorer on the left and a code editor on the right. The code editor displays a Jupyter Notebook cell with the following data:

	host	method	endpoint	protocol	response_code	content_size	dates	day	hour	minute	sec
1670441	ix-cin1-28.ix.netcom.com	GET	/shuttle/missions/sts-64/sts-64-patch-small.gif	HTTP/1.0	200	15980.0	1995-07-24 18:22:59	24	18	22	59
387438	tedspmac.dl.ac.uk	GET	/cgi-bin/imagemap/countdown?101,144	HTTP/1.0	302	96.0	1995-07-06 03:57:14	6	3	57	14
938353	192.87.128.134	GET	/ftbin/cdt_clock.pl	HTTP/1.0	200	752.0	1995-07-13 09:50:54	13	9	50	54
442487	www.npd.com	GET	/images/ksclgostsmall.gif	HTTP/1.0	304	0.0	1995-07-06 15:48:49	6	15	48	49
712273	165.113.187.62	GET	/history/apollo/apollo-13/apollo-13.html	HTTP/1.0	200	18114.0	1995-07-10 22:59:48	10	22	59	48
...	...	...	...	...	...	...	...	...	...	...	...
1304101	131.182.120.194	GET	/elv/ATLAS_CENTAUR/acsun.jpg	HTTP/1.0	200	12413.0	1995-07-18 13:44:32	18	13	44	32
418278	renata.atsc.allied.com	GET	/history/apollo/apollo-1/apollo-1-patch-small.gif	HTTP/1.0	200	16979.0	1995-07-06 12:16:41	6	12	16	41
1022449	kuts10p02.cc.ukans.edu	GET	/shuttle/countdown/count70.gif	HTTP/1.0	200	46573.0	1995-07-14 00:08:02	14	0	8	2
711609	tdosaj.bvlvigs.net	GET	/history/history.html	HTTP/1.0	200	1602.0	1995-07-10 22:44:37	10	22	44	37
1571316	200.26.8.6	GET	/images/ksclgost-medium.gif	HTTP/1.0	200	5866.0	1995-07-22 19:41:29	22	19	41	29

13818 rows x 11 columns

The cleaned data is now ready for transformation – normalization for numerical data and One Hot Encode from categorical data , we also discretize continuous attributes.

Host and endpoint have too many unique values and hence are not considered as features to be fed to the clustering algorithm,

The date-time column is removed too as information has been extracted from it into hour , min, sec ,day

**Continuous values:** Day , Hour, Min ,sec

These are discretized into bins of 5

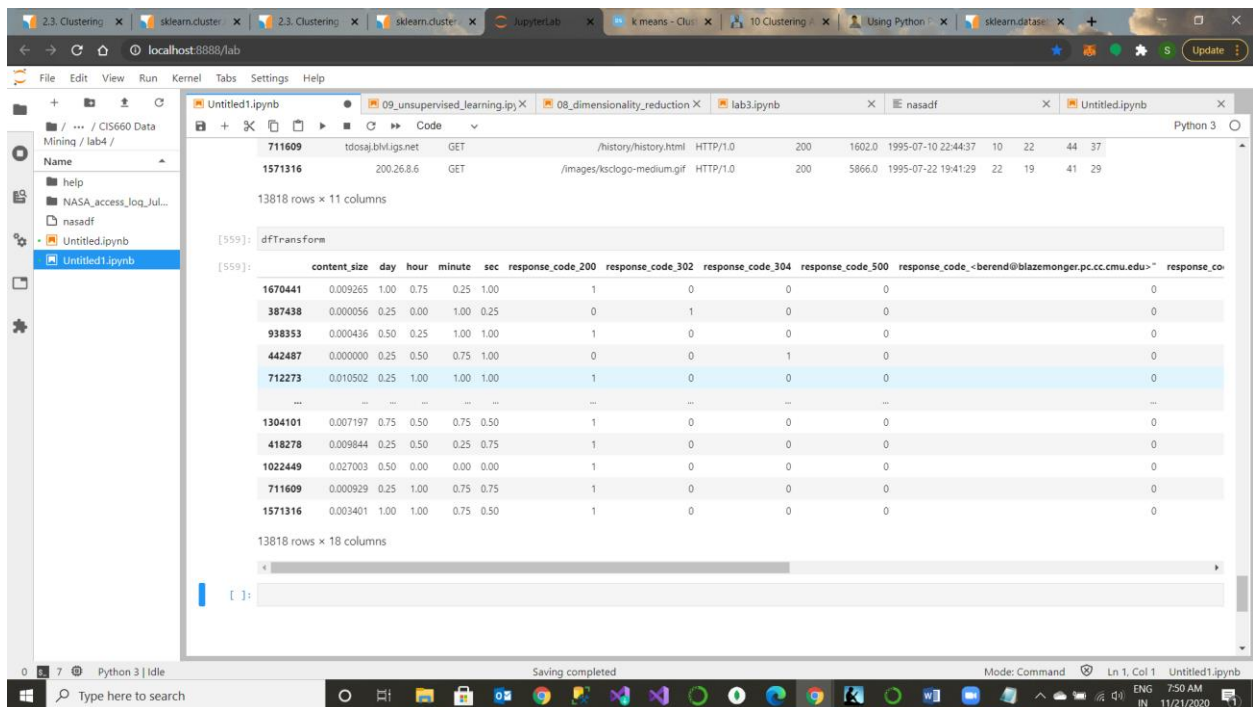
**Numerical Values:** Day , Hour, Min ,sec, content\_size

These values are normalized using the MinMax Scaler.

**Categorical values:** method , protocol ,response\_code

These values are one hot encoded

After data pre-processing the final training data looks as follows:

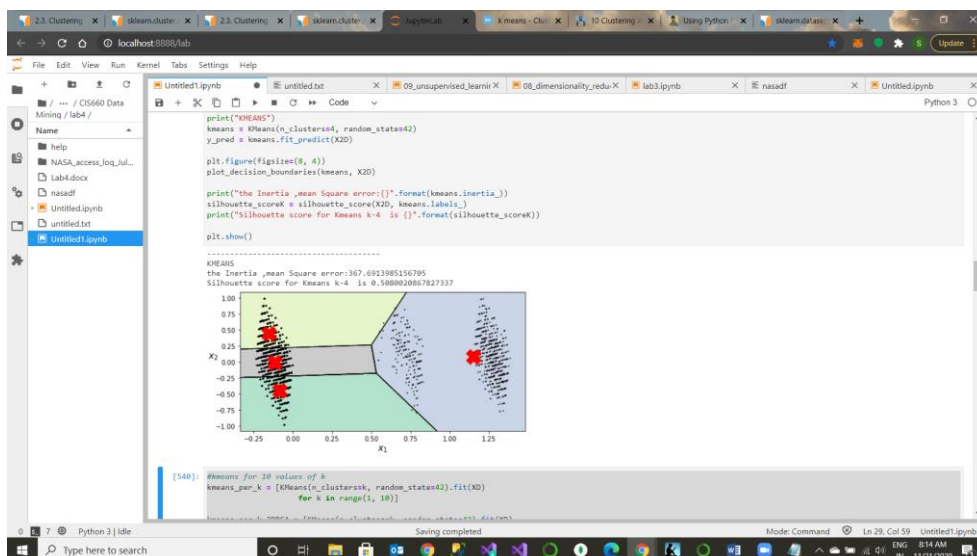


## PART-2 1 Data Analytic Experiment with Two different Clustering of Your Choice.

### K-Means clustering Algorithm

KMeans plot for k=4

the Inertia ,mean Square error:367.6913985156705  
Silhouette score for Kmeans k-4 is 0.5080020867827337



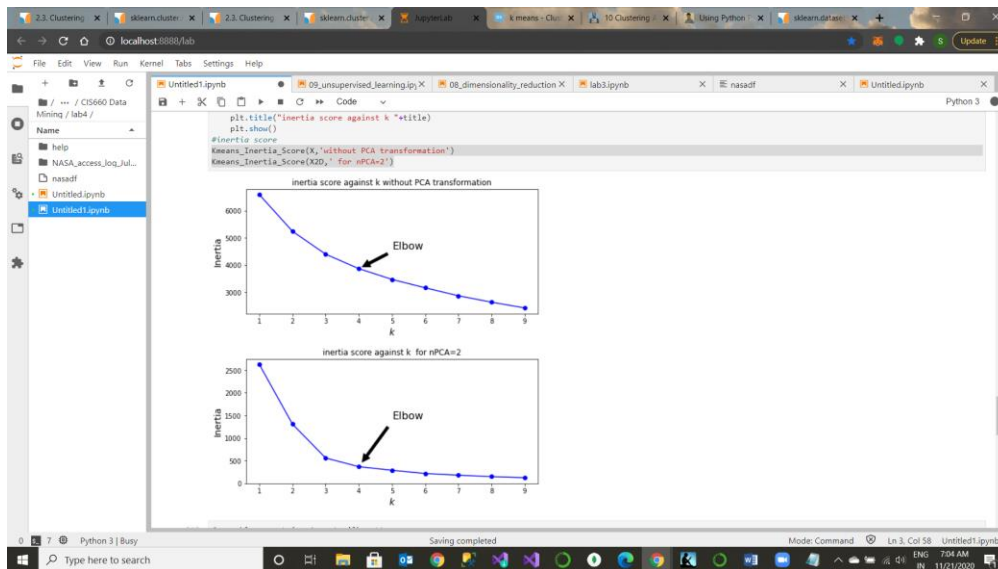
## EXTRA CREDIT

We consider the different values of  $k$  from 1 to 10 to train the kmeans algorithm on.

The silhouette score and the Mean Square Score for each of the  $k$  values is found and compared on a graph below.

The above is done on the training dataset as it is and also by applying PCA reduction tools with  $nPCA=2$

### Comparison of Inertia(MSE) scores for $k=1$ to 10 with PCA and non-PCA transformed dataset:

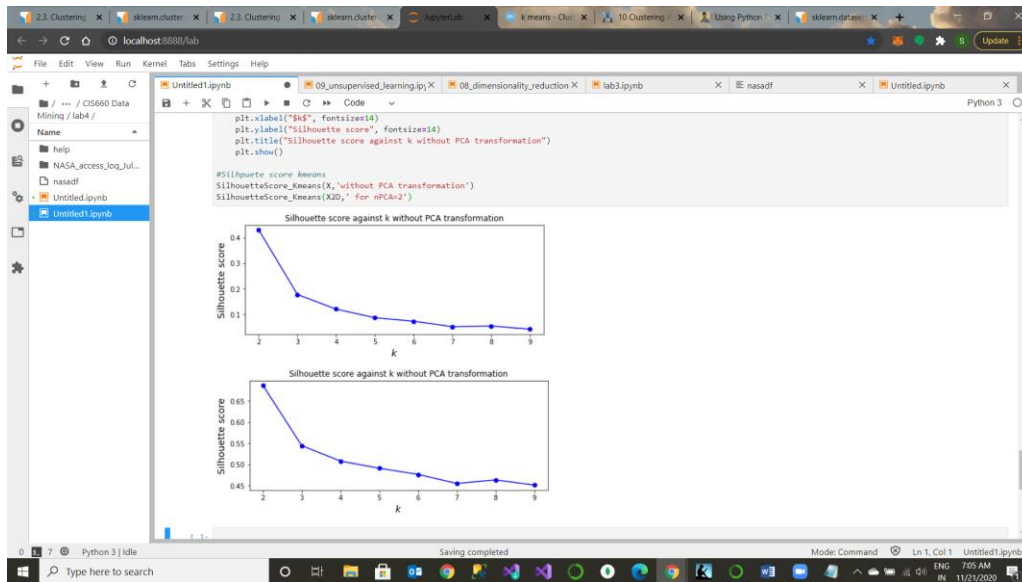


As you can see, there is an elbow at  $k=4$ , which means that less clusters than that would be bad, and more clusters would not help much and might cut clusters in half. So  $k=4$  is a pretty good choice.

we cannot simply take the value of  $k$  that minimizes the inertia, since it keeps getting lower as we increase  $k$ . Indeed, the more clusters there are, the closer each instance will be to its closest centroid, and therefore the lower the inertia will be.

### Comparison of Silhouette scores for $k=1$ to 10 with PCA and non-PCA transformed dataset:

As you can see, this visualization is much richer than the previous one: in particular, although it confirms that  $k=4$  is a very good choice, but it also underlines the fact that  $k=5$  is quite good as well.



Thus  $k=4, 5$  is the best choice for the given dataset when trained on KMEANS clustering algorithm.

## **DBSCAN**

We train two DBSCAN clustering algorithms one with  $ep=0.05$  and the other with  $ep=0.2$

Silhouette score for DBSCAN with  $eps=0.05$  is 0.7677809184623039

Silhouette score for DBSCAN with  $eps=0.05$  is 0.21553663221320152

And

Silhouette score for DBSCAN with  $eps=0.2$  is -0.10534056179625612

Silhouette score for DBSCAN with  $eps=0.2$  is 0.21553663221320152

Thus we understand that DBSCAN with  $eps=0.05$  is better than  $eps=0.2$

```

dbscan = DBSCAN(eps=0.05, min_samples=5)
dbscan.fit(X)

silhouette_scores_DBSCAN = silhouette_score(X, dbscan.labels_)
print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scores_DBSCAN))

dbscan.fit(X2D)

silhouette_scores_DBSCAN = silhouette_score(X2D, dbscan.labels_)
print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scores_DBSCAN))

DBSCAN
Silhouette score for DBSCAN with eps=0.05 is 0.7677809184623039
Silhouette score for DBSCAN with eps=0.05 is 0.21553663221320152

[555]: dbscan2 = DBSCAN(eps=0.2)
dbscan2.fit(X)
silhouette_scores_DBSCAN = silhouette_score(X, dbscan2.labels_)
print("Silhouette score for DBSCAN with eps=0.2 is {}".format(silhouette_scores_DBSCAN))
dbscan2.fit(X2D)
silhouette_scores_DBSCAN = silhouette_score(X2D, dbscan2.labels_)
print("Silhouette score for DBSCAN with eps=0.2 is {}".format(silhouette_scores_DBSCAN))

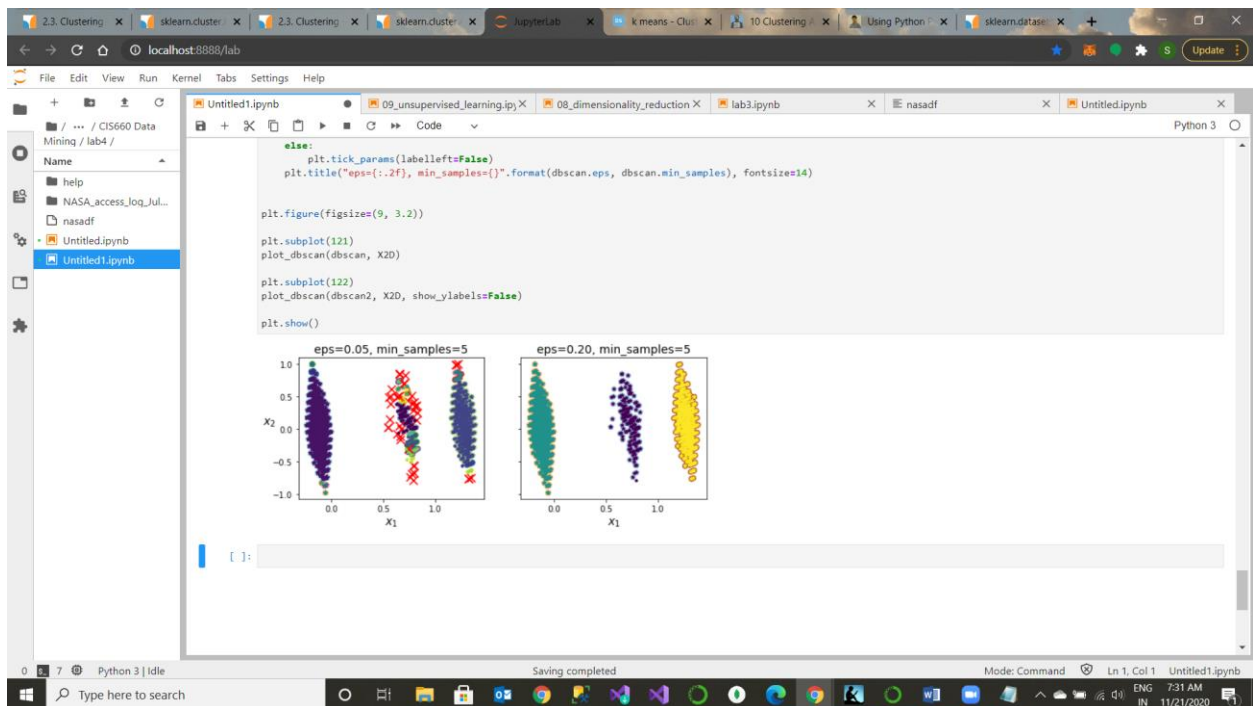
Silhouette score for DBSCAN with eps=0.2 is -0.10534056179625612
Silhouette score for DBSCAN with eps=0.05 is 0.21553663221320152

[556]: def plot_dbscan(dbscan, X, show_xlabel=True, show_ylabel=True):
    core_mask = np.zeros_like(dbscan.labels_, dtype=bool)
    core_mask[dbscan.core_sample_indices_] = True
    anomalies_mask = dbscan.labels_ == -1
    non_core_mask = ~(core_mask | anomalies_mask)

    cores = dbscan.components_
    anomalies = X[anomalies_mask]
    non_cores = X[non_core_mask]

    plt.scatter(cores[:, 0], cores[:, 1],
                c=dbscan.labels_[core_mask], markers='o', cmap="Paired")
    plt.scatter(cores[:, 0], cores[:, 1], markers='x', s=20, c=dbscan.labels_[core_mask])

```



## RESULTS:

- Kmeans is a least-squares optimization, whereas DBSCAN finds density-connected regions.

- Our experiment works best with DBSCAN as the data consists of dense regions separated by sparse regions and DBSCAN is good at clustering these kinds of data. Means works to minimize the least squares
- From our experiment we get that DBSCAN has better accuracy and performs better than Kmeans algorithm
- The silhouette score of DBSCAN is 0.75 whereas that of KMeans k=4 is 0.5
- We compared the inertias and silhouette scores of various k values and found that the best value of k=4 as there is an elbow at **k=4**, which means that less clusters than that would be bad, and more clusters would not help much and might cut clusters in half.
- But after finding the silhouette score we get that both k=4,5 are good options.
- We also see that KMeans performs better when PCA dimensionality reduction is done on the data before feeding it to KMeans algorithm. For our dataset nPCA=2 yields the best results.

#### **Source Code:**

```

from sklearn.cluster import KMeans
from sklearn.cluster import DBSCAN
import numpy as np
import re
import pandas as pd
import datetime

from sklearn import preprocessing
from sklearn.preprocessing import MinMaxScaler
import matplotlib.pyplot as plt
from sklearn.metrics import silhouette_score

N_file=open('NASA_access_log_Jul95/access_log_Jul95',encoding='latin1')

Lines=N_file.readlines()

df=pd.DataFrame([line.split() for line in Lines])

df.drop(columns=[i for i in range(10,56)],axis=1,inplace=True)

```



df

```
log_df=df.sample(14000)
```

```
log_df=log_df.dropna()
```

```
def parse_apache_time(s):
```

```
    month_map = {'Jan': 1, 'Feb': 2, 'Mar':3, 'Apr':4, 'May':5, 'Jun':6, 'Jul':7,  
                'Aug':8, 'Sep': 9, 'Oct':10, 'Nov': 11, 'Dec': 12}
```

```
    """ Convert Apache time format into a Python datetime object
```

```
    Args:
```

```
        s (str): date and time in Apache time format
```

```
    Returns:
```

```
        datetime: datetime object (ignore timezone for now)
```

```
    """
```

```
    return datetime.datetime(int(s[7:11]),
```

```
                              month_map[s[3:6]],
```

```
                              int(s[0:2]),
```

```
                              int(s[12:14]),
```

```
                              int(s[15:17]),
```

```
                              int(s[18:20]))
```

```
#removing rows with " in prtocol
```

```
log_df=log_df[log_df[8]!='']
```

```

#modifying the date columns

M=[l.replace('[',',') for l in log_df[3]]
dates=[parse_apache_time(s) for s in M]
log_df['dates']=dates
log_df.drop(columns=[1,2,3,4],inplace=True)

#renaming columns
log_df.rename(columns={0:'host', 5: 'method', 6: 'endpoint',
                        7:'protocol',
                        8:'response_code',
                        9:'content_size'},inplace=True)

log_df['method']=[l.replace(' ','') for l in log_df['method']]
#log_df['content_size']=[l.replace('\n','') for l in log_df['content_size']]
#removing empty protocol rows
log_df['protocol']=[l.replace(' ','') for l in log_df['protocol']]

#getting the dates
log_df['dates'] = pd.to_datetime(log_df['dates'],
format = '%Y-%m-%dT%H:%M:%SZ',
errors = 'coerce')

#log_df['year'] = log_df['dates'].dt.year
#log_df['month'] = log_df['dates'].dt.month
log_df['day'] = log_df['dates'].dt.day
log_df['hour'] = log_df['dates'].dt.hour
log_df['minute'] = log_df['dates'].dt.minute

```

```
log_df['sec'] = log_df['dates'].dt.second
```

```
log_df=log_df[log_df.content_size!='-']
```

```
log_df['content_size']=log_df['content_size'].astype('float64',copy=False)
```

```
#Data preprocessing
```

```
dfTransform=log_df[['method','protocol','response_code','content_size','day','hour','minute','sec']].copy()  
(
```

```
#normalization and discretization
```

```
minmax=MinMaxScaler()
```

```
dfTransform['day']=pd.qcut(log_df['day'], q=5,labels=[0,1,2,3,4])
```

```
dfTransform['day']=minmax.fit_transform(dfTransform[['day']])
```

```
dfTransform['hour']=pd.qcut(log_df['hour'], q=5,labels=[0,1,2,3,4])
```

```
dfTransform['hour']=minmax.fit_transform(dfTransform[['hour']])
```

```
dfTransform['minute']=pd.qcut(log_df['minute'], q=5,labels=[0,1,2,3,4])
```

```
dfTransform['minute']=minmax.fit_transform(dfTransform[['minute']])
```

```
dfTransform['sec']=pd.qcut(log_df['sec'], q=5,labels=[0,1,2,3,4])
```

```
dfTransform['sec']=minmax.fit_transform(dfTransform[['sec']])
```

```
dfTransform['content_size']=minmax.fit_transform(dfTransform[['content_size']])
```

```
#OneHotEncoding
```

```
OHE=pd.get_dummies(log_df['response_code'],prefix='response_code')
```

```
dfTransform=dfTransform.drop('response_code',axis=1)
```

```
dfTransform=dfTransform.join(OHE)
```

```
OHE=pd.get_dummies(log_df['method'],prefix='method')
```

```
dfTransform=dfTransform.drop('method',axis=1)
```

```
dfTransform=dfTransform.join(OHE)
```

```
OHE=pd.get_dummies(log_df['protocol'],prefix='protocol')
```

```
dfTransform=dfTransform.drop('protocol',axis=1)
```

```
dfTransform=dfTransform.join(OHE)
```

```
#splitting into train and validation
```

```
df_valid = dfTransform.sample(frac = 0.3, random_state = 42)
```

```
X = dfTransform.drop(df_valid.index)
```

```
from sklearn.decomposition import PCA
```

```
pca = PCA(n_components = 2)
```

```
X2D = pca.fit_transform(X)
```

```
pca = PCA(n_components = 5)
```

```
X5D = pca.fit_transform(X)
```

```
#plotting the data
```

```
def plot_data(X):
```

```

plt.plot(X[:, 0], X[:, 1], 'k.', markersize=2)

def plot_centroids(centroids, weights=None, circle_color='w', cross_color='r'):
    if weights is not None:
        centroids = centroids[weights > weights.max() / 10]
    plt.scatter(centroids[:, 0], centroids[:, 1],
                marker='o', s=30, linewidths=4,
                color=circle_color, zorder=10, alpha=0.9)
    plt.scatter(centroids[:, 0], centroids[:, 1],
                marker='x', s=50, linewidths=25,
                color=cross_color, zorder=11, alpha=1)

def plot_decision_boundaries(clusterer, X, resolution=1000, show_centroids=True,
                             show_xlabels=True, show_ylabels=True):
    mins = X.min(axis=0) - 0.1
    maxs = X.max(axis=0) + 0.1
    xx, yy = np.meshgrid(np.linspace(mins[0], maxs[0], resolution),
                          np.linspace(mins[1], maxs[1], resolution))
    Z = clusterer.predict(np.c_[xx.ravel(), yy.ravel()])
    Z = Z.reshape(xx.shape)

    plt.contourf(Z, extent=(mins[0], maxs[0], mins[1], maxs[1]),
                 cmap="Pastel2")
    plt.contour(Z, extent=(mins[0], maxs[0], mins[1], maxs[1]),
                linewidths=1, colors='k')
    plot_data(X)
    if show_centroids:
        plot_centroids(clusterer.cluster_centers_)

```

```

if show_xlabel:
    plt.xlabel("$x_1$", fontsize=14)
else:
    plt.tick_params(labelbottom=False)
if show_ylabel:
    plt.ylabel("$x_2$", fontsize=14, rotation=0)
else:
    plt.tick_params(labelleft=False)

k = 4
print("-----")
print("KMEANS")
kmeans = KMeans(n_clusters=4, random_state=42)
y_pred = kmeans.fit_predict(X2D)

plt.figure(figsize=(8, 4))
plot_decision_boundaries(kmeans, X2D)

print("the Inertia ,mean Square error:{}".format(kmeans.inertia_))
silhouette_scoreK = silhouette_score(X2D, kmeans.labels_)
print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scoreK))

plt.show()

#kmeans for 10 values of k
kmeans_per_k = [KMeans(n_clusters=k, random_state=42).fit(XD)
                 for k in range(1, 10)]

kmeans_per_k_2DPCA = [KMeans(n_clusters=k, random_state=42).fit(XD)

```

```

        for k in range(1, 10)]

def Kmeans_Inertia_Score(XD,title,kmeans_per_k):

    inertias = [model.inertia_ for model in kmeans_per_k]

    plt.figure(figsize=(8, 3.5))
    plt.plot(range(1, 10), inertias, "bo-")
    plt.xlabel("$k$", fontsize=14)
    plt.ylabel("Inertia", fontsize=14)
    plt.annotate('Elbow',
                  xy=(4, inertias[3]),
                  xytext=(0.55, 0.55),
                  textcoords='figure fraction',
                  fontsize=16,
                  arrowprops=dict(facecolor='black', shrink=0.1)
                  )

    plt.title("inertia score against k "+title)
    plt.show()

#inertia score
print("KMEANS Inertia- Mean square error score")

Kmeans_Inertia_Score(X,'without PCA transformation',kmeans_per_k)

Kmeans_Inertia_Score(X2D,' for nPCA=2',kmeans_per_k_2DPCA)

def SilhouetteScore(XD,title,labels,kmeans_per_k):

    silhouette_scoresO = [silhouette_score(XD, model.labels_)

```

```

        for model in kmeans_per_k[1:]

plt.figure(figsize=(8, 3))

plt.plot(range(2, 10), silhouette_scoresO, "bo-")

plt.xlabel("$k$", fontsize=14)

plt.ylabel("Silhouette score", fontsize=14)

plt.title("Silhouette score against k without PCA transformation")

plt.show()


#Silhouette score kmeans

SilhouetteScore(X,'without PCA transformation',kmeans.labels_,kmeans_per_k)

SilhouetteScore(X2D,' for nPCA=2',kmeans.labels_,kmeans_per_k_2DPCA)


#DBSCAN

from sklearn.cluster import DBSCAN

print("-----")

print("DBSCAN ")

dbscan = DBSCAN(eps=0.05, min_samples=5)

dbscan.fit(X)


silhouette_scores_DBSCAN = silhouette_score(X, dbscan.labels_)

print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scores_DBSCAN))


dbscan.fit(X2D)


silhouette_scores_DBSCAN = silhouette_score(X2D, dbscan.labels_)

print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scores_DBSCAN))


dbscan2 = DBSCAN(eps=0.2)

dbscan2.fit(X)

```



```

silhouette_scores_DBSCAN = silhouette_score(X, dbscan.labels_)
print("Silhouette score for DBSCAN with eps=0.2 is {}".format(silhouette_scores_DBSCAN))

dbscan2.fit(X2D)

silhouette_scores_DBSCAN = silhouette_score(X2D, dbscan.labels_)
print("Silhouette score for DBSCAN with eps=0.05 is {}".format(silhouette_scores_DBSCAN))

def plot_dbscan(dbscan, X, show_xlabels=True, show_ylabels=True):
    core_mask = np.zeros_like(dbscan.labels_, dtype=bool)
    core_mask[dbscan.core_sample_indices_] = True
    anomalies_mask = dbscan.labels_ == -1
    non_core_mask = ~(core_mask | anomalies_mask)

    cores = dbscan.components_
    anomalies = X[anomalies_mask]
    non_cores = X[non_core_mask]

    plt.scatter(cores[:, 0], cores[:, 1],
                c=dbscan.labels_[core_mask], marker='o', cmap="Paired")
    plt.scatter(cores[:, 0], cores[:, 1], marker='*', s=20, c=dbscan.labels_[core_mask])
    plt.scatter(anomalies[:, 0], anomalies[:, 1],
                c="r", marker="x", s=100)
    plt.scatter(non_cores[:, 0], non_cores[:, 1], c=dbscan.labels_[non_core_mask], marker=".")
    if show_xlabels:
        plt.xlabel("$x_1$", fontsize=14)
    else:
        plt.tick_params(labelbottom=False)
    if show_ylabels:
        plt.ylabel("$x_2$", fontsize=14, rotation=0)
    else:

```

```
plt.tick_params(labelleft=False)
plt.title("eps={:.2f}, min_samples={}".format(dbscan.eps, dbscan.min_samples), fontsize=14)
plt.figure(figsize=(9, 3.2))
plt.subplot(121)
plot_dbscan(dbscan, X2D)
plt.subplot(122)
plot_dbscan(dbscan2, X2D, show_ylabels=False)
plt.show()
```